

Implementation of SGDRegressor from Scratch:

```
In [ ]: #Using the Boston data set from the Sklearn Library :
```

```
In [29]: from sklearn.datasets import load_boston
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import SGDRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.cross_validation import train_test_split
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

boston = load_boston()
```

```
In [30]: #Loading the data from the boston dataset:
data = boston.data #featured data
target = boston.target #variable values

boston_df = pd.DataFrame(data)
X = boston_df
y = target
```

```
In [31]: # Standardize the data (mean=0, std=1) using training data
X_scaler = StandardScaler().fit(X)
standardized_X = X_scaler.transform(X)
```

Implementing the SGD Model :

```
In [32]: # The below function will compute the cost for each point:
def cal_cost(theta,X,y):
    m = len(y)

    predictions = X.dot(theta)

    cost = (1/2*m) * np.sum(np.square(predictions-y))

    return cost
```

```
In [227]: #The below method will compute the optimal weights and cost :
def stochastic_gradient_descent(X,y,theta,learning_rate=0.01,iterations=10):

    m = len(y) #Length of the data set
    cost_value = np.zeros(iterations)

    for it in range(iterations):
        cost = 0.0

        for i in range(m):
            rand_ind = np.random.randint(0,m)
            X_i = X[rand_ind,:].reshape(1,X.shape[1])
            y_i = y[rand_ind].reshape(1,1)
            prediction = np.dot(X_i,theta)

            theta = theta - (2/m)*learning_rate*( X_i.T.dot((prediction - y_i)))
            cost += cal_cost(theta,X_i,y_i)

        cost_value[it] = cost

    return theta, cost_value
```

```
In [228]: lr = 0.2 #Learning_rate
n_iter = 100 #no. of iterations

theta = np.random.randn(14,1)

X_b = np.c_[np.ones((len(standardized_X),1)),standardized_X] #adding the bias weight

theta_updated,cost_history = stochastic_gradient_descent(X_b,y,theta,lr,n_iter) #
```

```
In [229]: print('Intercept Term(bias term) : {:.3f}\n'.format(theta_updated[0][0]))

print('*'*100)
print('Predicted Weights(without bias term :)')
weight_vector = theta_updated[1:]
print(weight_vector)
```

Intercept Term(bias term) : 22.555

Predicted Weights(without bias term :)

```
[[-0.9638304 ]
 [ 0.9149475 ]
 [-0.01600848]
 [ 0.56743312]
 [-2.08737184]
 [ 2.80564371]
 [ 0.1813708 ]
 [-3.09703528]
 [ 2.45499635]
 [-1.75823064]
 [-2.11377119]
 [ 0.80181488]
 [-3.9207029 ]]
```

```
In [230]: y_predicted = X_b.dot(theta_updated)

y_predicted = y_predicted.ravel()
error = y - y_predicted
mse_sgd = np.mean(error**2)
print('Mean Squared Error using the predicted weights : ' , mse_sgd)
```

Mean Squared Error using the predicted weights : 21.991174425443887

Sklearn's SGDRegressor :

```
In [231]: #Computing the intercept and weight coffecients using the sklearn library:
```

```
lm = SGDRegressor(n_iter = 100 , penalty=None , loss='squared_loss' )
lm.fit(standardized_X , y)
```

```
Out[231]: SGDRegressor(alpha=0.0001, average=False, epsilon=0.1, eta0=0.01,
    fit_intercept=True, l1_ratio=0.15, learning_rate='invscaling',
    loss='squared_loss', max_iter=None, n_iter=100, penalty=None,
    power_t=0.25, random_state=None, shuffle=True, tol=None, verbose=0,
    warm_start=False)
```

```
In [232]: #Intercept Calculation:
print('Intercept term :', lm.intercept_)
```

Intercept term : [22.53564532]

```
In [233]: #Coffecient Calculation :
print('Weight vector :\n', lm.coef_.reshape(-1, 1))
```

```
Weight vector :
[[-0.91026934]
 [ 1.05890443]
 [ 0.11355234]
 [ 0.67833332]
 [-2.04260332]
 [ 2.65719259]
 [ 0.02343965]
 [-3.12238977]
 [ 2.55149977]
 [-1.9320931 ]
 [-2.05428794]
 [ 0.84663064]
 [-3.73052709]]
```

```
In [234]: #Predicting the target values for standardised data :
```

```
y_predict = lm.predict(standardized_X)
y_predict.shape
```

Out[234]: (506,)

```
In [235]: #mean square error
```

```
mse_sklearn = np.mean((y - y_predict)**2)
print('MSE:', mse_sklearn)
```

MSE: 21.90851533259455

Comparing our SGD model to Sklearn's SGDRegressor :

Bias Term (Intercept Term) :

```
In [236]: intercept = np.hstack([theta_updated[0][0], lm.intercept_])
intercept
```

Out[236]: array([22.55490109, 22.53564532])

Weight Vector :

```
► In [237]: weight = np.hstack([weight_vector , lm.coef_.reshape(-1 ,1)])  
weight
```

```
Out[237]: array([[ -0.9638304 , -0.91026934],  
                [  0.9149475 ,  1.05890443],  
                [ -0.01600848,  0.11355234],  
                [  0.56743312,  0.67833332],  
                [ -2.08737184, -2.04260332],  
                [  2.80564371,  2.65719259],  
                [  0.1813708 ,  0.02343965],  
                [ -3.09703528, -3.12238977],  
                [  2.45499635,  2.55149977],  
                [ -1.75823064, -1.9320931 ],  
                [ -2.11377119, -2.05428794],  
                [  0.80181488,  0.84663064],  
                [ -3.9207029 , -3.73052709]])
```

MSE Comparison:

```
In [238]: mse_array = np.hstack([mse_sgd , mse_sklearn])  
mse_array
```

```
Out[238]: array([21.99117443, 21.90851533])
```

CONCLUSION:

1) We got same approximately same values for the weight vector , intercept term and Mean Squared Error values in case of SGD from scratch and Sklearn SGD model.

2) So our model is working fine with learning rate(eat0) = 0.2 and n_iter = 100.

```
In [ ]:
```